

A reference-cohort summary estimator for genetic and gene–environment variance components

SUMMIT technical report

14 August 2026

Abstract

GENIE estimates additive genetic (G), gene–environment (G×E), noise–environment (N×E), and residual variance components from one individual-level cohort. Rewriting GENIE’s right-hand side in terms of marginal GWAS and GWIS scores is exact when the kernel traces and the scores come from the same individuals. It is not, by itself, a reusable reference method. This report derives the additional moments needed when G and G×E kernel quantities are estimated in an independent cohort containing all individuals with measured environment, while trait-specific marginal scores may use smaller and different samples. The key is to transfer the same-individual and different-individual parts of each genetic kernel-product trace separately. For G×E, the same-individual term is a non-universal fourth moment and cannot be replaced by the additive-only constant. It can be estimated during the existing randomized, in-memory reference pass with an order-two probe U-statistic and $O(N_R K)$ additional memory, independent of the number of probes. Study-specific N×E moments remain exact low-dimensional summaries. We also audit covariate handling in the released GENIE code and distinguish two coherent feature definitions from the incoherent mixture of raw projected kernels and hard-coded traces.

1 Notation and generative assumptions

We follow the notation of Kerin and Marchini [1]. Cohort S is the trait study and cohort R is the reusable reference. They share an ordered set of M variants and the same annotation weights $A_{jk} \geq 0$, $k = 1, \dots, K$, with $M_k = \sum_j A_{jk}$. Their individuals need not be the same. Let $\mathbf{C}_S$ contain an intercept, the environment, and the other fixed effects used by the trait analysis. Define

$$\mathbf{P}_S = \mathbf{I} - \mathbf{C}_S(\mathbf{C}_S^\top \mathbf{C}_S)^+ \mathbf{C}_S^\top, \quad r_S = \text{tr}(\mathbf{P}_S) = N_S - \text{rank}(\mathbf{C}_S), \quad (1)$$

where $(\cdot)^+$ is the Moore–Penrose inverse. The same definitions apply to R .

Let $\mathbf{G}_S$ be a centered, pre-scaled genotype matrix and let $\mathbf{e}_S$ be a centered environment measured in standard-deviation units. The default one-environment GENIE generative model, before fixed-effect projection, is

$$\tilde{\mathbf{y}}_S = \mathbf{G}_S \tilde{\boldsymbol{\beta}} + \text{diag}(\mathbf{e}_S) \mathbf{G}_S \tilde{\boldsymbol{\alpha}} + \text{diag}(\mathbf{e}_S) \boldsymbol{\delta} + \mathbf{C}_S \boldsymbol{\gamma} + \boldsymbol{\varepsilon}. \quad (2)$$

The published framework allows several environments; the present derivation applies to one environment at a time. GENIE uses the naturally scaled projected columns $\mathbf{P}_S \mathbf{G}_S$ and $\mathbf{P}_S \text{diag}(\mathbf{e}_S) \mathbf{G}_S$.

SUMMIT's default instead restandardizes those columns after projection:

$$\begin{aligned} \mathbf{U}_S &= \mathbf{P}_S \mathbf{G}_S, & \mathbf{X}_{S,:,j} &= \mathbf{U}_{S,:,j} \sqrt{\frac{r_S}{\mathbf{U}_{S,:,j}^\top \mathbf{U}_{S,:,j}}}, \\ \mathbf{V}_S &= \mathbf{P}_S \text{diag}(\mathbf{e}_S) \mathbf{G}_S, & \mathbf{W}_{S,:,j} &= \mathbf{V}_{S,:,j} \sqrt{\frac{r_S}{\mathbf{V}_{S,:,j}^\top \mathbf{V}_{S,:,j}}}. \end{aligned} \quad (3)$$

Thus every retained column of $\mathbf{X}_S$ and $\mathbf{W}_S$ has squared norm r_S . This is the G×E analogue of SUMMIT's covariate-residualized, restandardized additive feature definition [3]. The two choices define different random-effect scalings, but the method-of-moments algebra below is the same when all traces and scores use one choice coherently. The external implementation uses the standardized choice. Define

$$\begin{aligned} \mathbf{K}_{Sk}^G &= \frac{\mathbf{X}_S \text{diag}(\mathbf{A}_{:,k}) \mathbf{X}_S^\top}{M_k}, & \mathbf{K}_{Sk}^I &= \frac{\mathbf{W}_S \text{diag}(\mathbf{A}_{:,k}) \mathbf{W}_S^\top}{M_k}, \\ \mathbf{D}_S &= \mathbf{P}_S \text{diag}(\mathbf{e}_S^2) \mathbf{P}_S, & \mathbf{K}_S^0 &= \mathbf{P}_S. \end{aligned} \quad (4)$$

The superscript I denotes interaction and avoids confusing an environment vector with an expectation.

The residualized phenotype is generated under the second-moment model

$$\mathbf{y}_S = \sum_{k=1}^K \mathbf{X}_S \boldsymbol{\beta}_k + \sum_{k=1}^K \mathbf{W}_S \boldsymbol{\alpha}_k + \mathbf{P}_S \text{diag}(\mathbf{e}_S) \boldsymbol{\delta} + \mathbf{P}_S \boldsymbol{\varepsilon}, \quad (5)$$

with mutually uncorrelated, mean-zero effects

$$\begin{aligned} \text{Cov}(\boldsymbol{\beta}_k) &= \frac{\sigma_{Gk}^2}{M_k} \text{diag}(\mathbf{A}_{:,k}), & \text{Cov}(\boldsymbol{\alpha}_k) &= \frac{\sigma_{Ik}^2}{M_k} \text{diag}(\mathbf{A}_{:,k}), \\ \text{Cov}(\boldsymbol{\delta}) &= \sigma_D^2 \mathbf{I}, & \text{Cov}(\boldsymbol{\varepsilon}) &= \sigma_0^2 \mathbf{I}. \end{aligned}$$

For overlapping annotations this is understood as a linear covariance model; the coefficients need not each be nonnegative. Conditional on the design,

$$\boldsymbol{\Sigma}_S = \text{Cov}(\mathbf{y}_S \mid \mathbf{X}_S, \mathbf{W}_S, \mathbf{e}_S, \mathbf{C}_S) = \sum_k \sigma_{Gk}^2 \mathbf{K}_{Sk}^G + \sum_k \sigma_{Ik}^2 \mathbf{K}_{Sk}^I + \sigma_D^2 \mathbf{D}_S + \sigma_0^2 \mathbf{P}_S. \quad (6)$$

No Gaussian assumption on the effects is required for the moment equations.

2 GENIE as an in-sample method-of-moments estimator

Index the $2K + 2$ kernels in Eq. (6) by $\mathcal{K}_S = (\mathbf{K}_{S1}^G, \dots, \mathbf{K}_{SK}^G, \mathbf{K}_{S1}^I, \dots, \mathbf{K}_{SK}^I, \mathbf{D}_S, \mathbf{P}_S)$ and collect their coefficients in $\boldsymbol{\theta}$. GENIE minimizes the Frobenius loss

$$\left\| \mathbf{y}_S \mathbf{y}_S^\top - \sum_a \theta_a \mathcal{K}_{Sa} \right\|_F^2. \quad (7)$$

Differentiation gives the normal equations

$$\mathbf{T}_S \hat{\boldsymbol{\theta}} = \mathbf{q}_S, \quad (\mathbf{T}_S)_{ab} = \text{tr}(\mathcal{K}_{Sa} \mathcal{K}_{Sb}), \quad (\mathbf{q}_S)_a = \mathbf{y}_S^\top \mathcal{K}_{Sa} \mathbf{y}_S. \quad (8)$$

This is precisely Supplemental Note S1 of GENIE after writing the projected kernels explicitly [2]. GENIE estimates large traces by Hutchinson random projection: for independent isotropic probes $\mathbf{z}_v$,

$$\text{tr}(\mathcal{K}_a \mathcal{K}_b) = \mathbb{E}_v[(\mathcal{K}_a \mathbf{z}_v)^\top (\mathcal{K}_b \mathbf{z}_v)]. \quad (9)$$

For a genetic family $F \in \{G, I\}$, define the marginal HE-scale score

$$s_{Sj}^F = \frac{\mathbf{F}_{S,:,j}^\top \mathbf{y}_S}{\sqrt{r_S}}. \quad (10)$$

It follows exactly that

$$\mathbf{y}_S^\top \mathbf{K}_{Sk}^F \mathbf{y}_S = \frac{r_S}{M_k} \sum_j A_{jk} (s_{Sj}^F)^2. \quad (11)$$

The required statistics are marginal cross-products after the common fixed-effect projection. A usual conditional interaction Wald statistic is not interchangeable with Eq. (10) unless its numerator and scale are transformed back to this marginal score.

Let

$$L_{S,jk}^{FH} = \sum_m A_{mk} \left(\frac{\mathbf{F}_{S,:,j}^\top \mathbf{H}_{S,:,m}}{r_S} \right)^2. \quad (12)$$

Then the genetic block of Eq. (8) is also exact:

$$\text{tr}(\mathbf{K}_{Sk}^F \mathbf{K}_{S\ell}^H) = \frac{r_S^2}{M_k M_\ell} \sum_j A_{jk} L_{S,j\ell}^{FH}. \quad (13)$$

Equations (11) and (13) explain the existing matched-cohort summary implementation. They do *not* justify inserting an independently estimated L_R into Eq. (13): the finite-cohort trace has both same-person fourth moments and different-person LD moments, which scale differently with cohort size.

3 A genuinely reusable reference estimator

3.1 Population transport assumptions

The reference construction requires the following assumptions.

1. Cohorts R and S are independent draws from the same target population, with the same variant/allele and annotation definitions.
2. After applying the declared fixed-effect convention and per-cohort restandardization, the joint second and fourth moments of (X, W, E) are transportable. Fixed-effect dimension is small relative to cohort size, so the projection's dependence and leverage perturb the exchangeable-individual moments only at lower order. Large leverage differences or population/LD/environment shift violate this assumption.
3. Trait missingness may change N_S but is ignorable conditional on the modeled variables. Genotypes, environment, and covariates used to form the marginal scores obey the same feature convention as the reference.

These are reference-panel assumptions, not algebraic identities. If the joint G–E distribution differs, the estimator converges to a reference-weighted pseudo-parameter, just as additive SUMMIT with mismatched LD has target $(\mathbf{A}_R)^{-1} \mathbf{A}_S \mathbf{h}^2$ [3].

3.2 Same-person and different-person trace moments

For any genetic kernels $a = (F, k)$ and $b = (H, \ell)$, write

$$\begin{aligned} D_{ab}^C &= \sum_{i=1}^{N_C} K_a^C(i, i) K_b^C(i, i), \\ O_{ab}^C &= \sum_{i \neq i'} K_a^C(i, i') K_b^C(i, i'), \quad T_{ab}^C = D_{ab}^C + O_{ab}^C. \end{aligned} \quad (14)$$

The split is on the individual axis, not an arbitrary orthonormal basis for the residual subspace. Residual rank r_C controls feature and score normalization, but it is not the number of exchangeable people. Define population moments

$$d_{ab} = \mathbb{E}[K_a(i, i) K_b(i, i)], \quad o_{ab} = \mathbb{E}[K_a(i, i') K_b(i, i')], \quad i \neq i'. \quad (15)$$

Then

$$\mathbb{E}[T_{ab}^S] = N_S d_{ab} + N_S(N_S - 1) o_{ab}. \quad (16)$$

The reference estimators are

$$\hat{d}_{ab} = \frac{\hat{D}_{ab}^R}{N_R}, \quad \hat{o}_{ab} = \frac{\hat{T}_{ab}^R - \hat{D}_{ab}^R}{N_R(N_R - 1)}, \quad (17)$$

and therefore

$$\boxed{\hat{T}_{ab}^{S|R} = \frac{N_S}{N_R} \hat{D}_{ab}^R + \frac{N_S(N_S - 1)}{N_R(N_R - 1)} (\hat{T}_{ab}^R - \hat{D}_{ab}^R).} \quad (18)$$

When $N_S = N_R$, Eq. (18) returns the original reference trace exactly. For additive standardized genotypes, D^R is close to the familiar N_R -order baseline. For $G \times E$, it contains quantities such as $\mathbb{E}[X^2 E^2 X'^2 E'^2]$ and is not universal. This is the precise reason that blind squared-rank scaling, or importing the additive leading term, fails. There is no assumption-free way to recover a study's conditional trace from an independent cohort. Equation (18) is the reference-panel estimator under the population-transport assumptions above; without them the target is necessarily reference weighted.

3.3 Estimating the fourth-moment term without a second genotype pass

For feature family $a = (F, k)$ and probe v , the reference trace pass already forms

$$S_{aiv} = \sum_j F_{R,ij} \sqrt{A_{jk}} z_{jv}, \quad \mathbb{E}_z[S_{aiv}^2] = M_k K_a^R(i, i). \quad (19)$$

For $B \geq 2$, an unbiased order-two probe U-statistic is

$$\hat{D}_{ab}^R = \frac{1}{B(B-1)M_k M_\ell} \sum_i \left[\left(\sum_v S_{aiv}^2 \right) \left(\sum_u S_{biu}^2 \right) - \sum_v S_{aiv}^2 S_{biv}^2 \right]. \quad (20)$$

The subtraction excludes identical probes; different probe indices are independent even when annotations overlap and $G/G \times E$ sketches share the same underlying probe. An implementation needs an $N_R \times 2K$ matrix of running square sums and one $2K \times 2K$ accumulator. Temporary conversion is chunked over probes. Memory is $16N_R K + O(K^2)$ bytes in float64 and is independent of B ; at $N_R = 291,273, K = 1$ it is about 4.4 MiB. No source sketch, block sketch, or cache is written to disk.

3.4 N×E and residual moments

N×E is study specific and cheap; it should not be borrowed from R . Let $\mathbf{Q}_S$ be an orthonormal basis for the fixed effects and $\mathbf{d}_S = \mathbf{e}_S^2$. With $\mathbf{D}_S = \mathbf{P}_S \text{diag}(\mathbf{d}_S) \mathbf{P}_S$,

$$\text{tr}(\mathbf{D}_S) = \sum_i d_{Si} - \text{tr}(\mathbf{Q}_S^\top \text{diag}(\mathbf{d}_S) \mathbf{Q}_S), \quad (21)$$

$$\text{tr}(\mathbf{D}_S^2) = \sum_i d_{Si}^2 - 2 \text{tr}(\mathbf{Q}_S^\top \text{diag}(\mathbf{d}_S^2) \mathbf{Q}_S) + \text{tr}\{(\mathbf{Q}_S^\top \text{diag}(\mathbf{d}_S) \mathbf{Q}_S)^2\}. \quad (22)$$

The phenotype supplies

$$q_D = \mathbf{y}_S^\top \mathbf{D}_S \mathbf{y}_S = \|\mathbf{e}_S \odot \mathbf{y}_S\|^2, \quad q_0 = \mathbf{y}_S^\top \mathbf{y}_S = r_S, \quad (23)$$

on SUMMIT's normalized residual scale. Cross-traces with genetic kernels are also exact study summaries. The same pass that computes marginal GWAS/GWIS scores forms

$$d_{Sj}^F = \frac{\mathbf{F}_{S,:,j}^\top \text{diag}(\mathbf{e}_S^2) \mathbf{F}_{S,:,j}}{r_S}, \quad \text{tr}(\mathbf{K}_{Sk}^F \mathbf{D}_S) = \frac{r_S}{M_k} \sum_j A_{jk} d_{Sj}^F, \quad (24)$$

because projected feature columns satisfy $\mathbf{P}_S \mathbf{F}_{S,:,j} = \mathbf{F}_{S,:,j}$. This costs no extra genotype decode and avoids transporting an environment-sensitive fourth moment. Because Eq. (3) makes $\text{tr}(\mathbf{K}_a^S) = r_S$, genetic-residual cross-traces are exact.

Equations (11), (18), (21)–(24) fill every entry of the study normal equations. Solving them estimates the four classes ($G, G \times E, N \times E$, residual) using genotype/environment moments from all eligible reference individuals and trait-specific marginal scores from the study cohort.

3.5 The summary-data boundary

The reusable reference object contains the ordered variant/allele and annotation definitions, N_R, r_R , the randomized genetic trace matrix $\hat{\mathbf{T}}_R$, and the compact same-individual matrix $\hat{\mathbf{D}}_R$. These quantities are phenotype independent and are computed once in every individual with an observed environment and reference covariates. For each trait, the required study object is

$$\mathcal{S}_y = \{N_S, r_S, \mathbf{s}_S^G, \mathbf{s}_S^I, q_D, q_0, \text{tr}(\mathbf{D}_S), \text{tr}(\mathbf{D}_S^2), \text{tr}(\mathbf{K}_{Sk}^G \mathbf{D}_S), \text{tr}(\mathbf{K}_{Sk}^I \mathbf{D}_S)\}_{k=1}^K. \quad (25)$$

The two score vectors are marginal GWAS/GWIS cross-products. The remaining terms are scalars or $O(K)$ design summaries. Equation (24) lets the genetic-N×E terms be accumulated in the same streamed genotype pass that produces GWAS/GWIS, with no randomized trace calculation and no second decode. Once $\mathcal{S}_y$ is written, fitting uses no individual-level phenotype, genotype, or environment. Thus trait missingness affects the cheap trait summary but does not force reconstruction of the all-E G×E reference.

3.6 Uncertainty and identifiability

There are three uncertainty sources: trait-score sampling, finite reference sampling, and randomized-probe error. A SNP-block jackknife deletes the corresponding GWAS/GWIS score rows and reference LD-score rows. The default block-local approximation assumes cross-block LD is negligible and reuses the full-reference $\hat{\mathbf{D}}^R$ in each replicate. This omits a change of order the deleted SNP fraction in the diagonal term, so it is an explicit approximation whose stability should be checked when

blocks are few, large, or annotation-imbalanced. An exact deletion would require block-specific cross-sketch moments and is intentionally not the production default.

Increasing B reduces only probe error. It cannot repair reference mismatch or an unidentified model. In particular, if E^2 is constant after projection, $\mathbf{D}_S$ is proportional to $\mathbf{P}_S$ and $N \times E$ is aliased with residual noise. Collinear annotation kernels cause analogous rank failure.

4 Validation

Three independent checks were used.

1. A dense oracle forms all individual-level kernels and verifies that the matched summary equations reproduce every GENIE matrix, right-hand side, and trace entry to floating-point tolerance (maximum tested discrepancy below 2×10^{-13}).
2. Setting $N_S = N_R$ in the implemented transfer reproduces the complete reference genetic block exactly, irrespective of how T_R is split into D_R and O_R .
3. In a separate continuous- E Monte Carlo experiment, an independent $N_R = 2500$ reference and 60 independent $N_S = 350$ studies used $M = 60$ correlated Gaussian markers (AR(1) correlation 0.35), with intercept and E projected out. Median relative errors of Eq. (18) were 2.17% (G–G), 3.14% (G–G×E), and 2.23% (G×E–G×E). Blind squared-rank scaling gave 9.27%, 86.46%, and 23.51%, respectively. Genetic- $N \times E$ cross-traces were computed exactly in each study rather than transported. These are finite-reference errors, not claimed universal operating characteristics.

An end-to-end integration test also constructs one reference, removes a trait value to create a different study cohort, writes the study GWAS/GWIS moments, and fits the transferred equations.

5 Covariates and restandardization in released GENIE code

The equations in the GENIE paper are coherent: with $\mathbf{V} = \mathbf{I} - \mathbf{C}(\mathbf{C}^\top \mathbf{C})^{-1} \mathbf{C}^\top$, Supplemental Note S1 uses $\text{tr}(\mathbf{K}_k \mathbf{V} \mathbf{K}_\ell \mathbf{V})$, $\text{tr}(\mathbf{V} \mathbf{K}_k)$, $\mathbf{y}^\top \mathbf{V} \mathbf{K}_k \mathbf{V} \mathbf{y}$, and residual dimension $N - C$ [2]. The concerns below apply to the released code at commit 22bee9c, not to that formal derivation.

Ordinary inverse and column count. The implementation computes $\mathbf{Q} = \mathbf{W} \mathbf{t} \mathbf{W}.\text{inverse}()$ and sets the residual dimension to $N_{\text{indv_mask}} - N_{\text{cov}}$ [4]. This is correct only if the retained covariate columns are exactly full rank and numerically well conditioned. Environment, intercept, duplicated indicators, or a constant covariate can make that false. A rank-revealing QR/SVD and $N - \text{rank}(\mathbf{C})$ are required for the projector in Eq. (1).

Scaling before the analysis mask. The genotype code computes allele means over N_{indv} and applies the HWE scale $\{2p(1-p)\}^{-1/2}$ before the phenotype/environment mask and fixed-effect projection [5]. Consequently a retained, projected column generally does not have squared norm N_{mask} or $N_{\text{mask}} - \text{rank}(\mathbf{C})$. This is especially visible when trait missingness creates a much smaller analysis sample.

The important distinction: two coherent estimands and one incoherent hybrid. There are two valid choices.

1. *Raw projected kernel:* use $U = PX$ and $W = P(E \odot X)$ without post-projection scaling. This defines a valid model, but every trace, including $\text{tr}(K)$, must be computed from the realized columns and the random effect prior is tied to the pre-projection genotype scale.
2. *Restandardized projected kernel:* rescale every projected column as in Eq. (3). Then all genetic kernel traces equal r and summary moments have a stable, cohort-comparable per-residual-SD definition. This is SUMMIT’s default.

The released GENIE implementation initializes each additive genetic trace to `Nindv_mask`, while interaction/noise traces are estimated from the random projections, and then applies a projection subtraction [6]. That mixes the raw-kernel Gram and phenotype moments with a unit-norm additive trace assumption that the preceding HWE scaling/masking does not guarantee. The problem is not simply “GENIE fails to restandardize”: raw projected kernels are legitimate. The problem is using their realized Gram/RHS while hard-coding a trace from the alternative standardized-kernel convention.

Phenotype scaling. The default code standardizes the phenotype, projects covariates, and then standardizes the projected phenotype again using denominator $N_{\text{mask}} - 1$ [7], while the residual kernel dimension is declared $N_{\text{mask}} - N_{\text{cov}}$. Because the code retains the realized $\mathbf{y}^\top \mathbf{y}$ in the normal equations, phenotype rescaling alone is not necessarily a bias. It does, however, mean that reported components are not automatically on the r -normalized scale and should not be combined with hard-coded r traces without an explicit scale conversion.

These issues can be small when covariates are full rank, missingness is mild, and HWE scaling nearly matches the retained sample variance. They are not algebraically controlled in the released implementation. SUMMIT instead uses a rank-revealing orthonormal basis, records the numerical rank, applies the trait/reference mask before feature construction, and either restandardizes projected columns or computes their actual traces consistently.

6 Computational implications

The new population statistic does not change the dominant complexity. With B probes, K annotations, M variants, and streamed genotype blocks, the G/G×E reference remains linear in N_R , M , B , and K for $K = 1$ (and in the number of requested directional annotation products generally). The added Eq. (20) work is $O(N_R B K^2)$ but has a very small constant and no genotype decode or large matrix multiply; its persistent memory is $O(N_R K)$.

On tabla, a sealed full $N_R = 291,273$, $M = 454,207$, $B = 32$ EUR age reference completed its numerical work in 1,503.6 seconds; end-to-end measured elapsed time was 1,534.9 seconds (25.6 minutes), with 4.04 GiB peak RSS. Its feature, source, panel-preparation, and target phases took 560.0, 368.6, 0.20, and 471.4 seconds. The source tree, loaded native binary, and single OpenBLAS 0.3.34 runtime are hash-bound in the published reference manifest. Relative to the preceding internally threaded diagnostic, total elapsed time fell 20.9%, target time fell 51.9%, and repaired output columns fell from 384 to 3 without increasing memory.

A direct $B/32$ extrapolation is too pessimistic because the $B = 32$ products are skinny. In a dense-oracle benchmark with $N = 100,000$ and $M = 1,000$, increasing B from 32 to 512 increased native source time from 0.319 to 1.267 seconds (3.97-fold) and the relevant $2B$ target time from

0.458 to 2.396 seconds (5.23-fold), rather than 16-fold. An affine fit to the measured $B = 32, 128, 512$ phase timings, scaled by the sealed full-cohort $B = 32$ phases, gives a point estimate near 2.2 hours at $B = 1024$. A conservative current planning range is 2.2–2.8 hours to allow for repair and system-load variation. This remains an extrapolation, not a completed $B = 1024$ production timing, and should be replaced by a measured value before a large production campaign.

Probe tiling and block-width adaptation bound memory. For one annotation, the float32 $[XZ, WZ]$ source sketch costs $8N_R B$ bytes. The native opaque $[S, E \odot S]$ target panel costs $32N_R B$ bytes, while the protected non-decoded target operand costs at most another $16N_R B$ bytes *inside* the configured native workspace. At $N_R = 291,273$ and $B = 1024$, these terms are about 2.22, 8.89, and 4.44 GiB. The first two allocations plus the bounded 16-GiB native workspace (whose variant block shrinks as the protected operand grows) imply an expected peak of roughly 25–30 GiB, not 32 times the $B = 32$ RSS. The new population-transfer statistic itself adds only about 4.4 MiB for $K = 1$.

The numerical failures are not caused by extreme Rademacher probes: their entries are exactly ± 1 . Nor is float32 the arithmetic root cause. The decoder, feature moments, projection, source/target GEMMs, checksums, and repairs use float64; float32 controls only retained sketch buffers. Switching those buffers to float64 doubles their memory and reduces accumulation roundoff, but cannot prevent a vendor GEMM from sparsely changing a double input or output. Dropping extreme LD-score values would be unsafe here: an extreme Monte Carlo realization and a corrupt output are different events, and clipping squared traces creates downward bias.

Before the environment repair, *tabla* loaded two BLAS runtimes in one process: NumPy’s pip wheel bundled OpenBLAS 0.3.30/Haswell while the native extension used conda OpenBLAS 0.3.34/Zen. The environment has now been changed to conda NumPy 2.3.5, so NumPy and the extension both resolve to the same OpenBLAS 0.3.34 Zen library. This removes a concrete source of process-wide thread and runtime interference, but it did not remove the numerical fault. The full-cohort $B = 32$ diagnostic triggered 384 repaired output columns. More decisively, a production-shaped dense test with $N = 291,273$, $M = 1,207$, and $B = 32$ intermittently observed a maximum target error of 58.15 from the internally threaded vendor path, despite using binary64 arithmetic. Repeating the same test could also return only ordinary accumulation-order error, proving that the failure is intermittent and shape/thread dependent.

SUMMIT therefore no longer invokes an internally threaded OpenBLAS GEMM for these products. It fixes OpenBLAS at one thread per call and partitions disjoint output rows across SUMMIT-owned OpenMP workers. This was faster, not slower, in the production-shaped no-fault comparison (1.09 versus 1.33 seconds for the $2B$ target and 1.79 versus 2.28 seconds for $4B$). Across 20 repeated $2B/4B$ calls, the selective ABFT layer repaired 257 output columns and every final output remained within 2.283×10^{-6} of the independent dense float64 oracle. Thus the vendor/concurrency problem is still observable, but the returned result is protected without duplicating the full genotype operand. The ABFT path checks wide GEMMs, repairs only discrepant output columns with a non-vendor cache-tiled kernel, and falls back to a deterministic fresh decode only if the reconstructable genotype input changes.

7 Practical contract

The implemented external path is explicit. A reference generated with at least two probes records the compact population trace matrix in its manifest. Trait scoring with `--gxe-population-reference` permits a different retained sample, recomputes study feature standardization, and writes study N , r , q_D , q_0 , $\text{tr}(D)$, $\text{tr}(D^2)$, and exact genetic- D cross-traces (plus delete-block ver-

sions when requested). The fit requires the exact reference-manifest hash, identical ordered variants/alleles and annotations, the same named environment/covariate convention, and standardized kernels. It fails closed for GENIE/raw kernels because their genetic traces are not fixed by r . Existing matched-cohort behavior remains the default.

References

- [1] Kerin M, Marchini J. Estimation of gene–environment interaction heritability from large-scale biobanks. *American Journal of Human Genetics* (2024). doi:10.1016/j.ajhg.2024.05.015.
- [2] Kerin M, Marchini J. GENIE Supplemental Notes, especially Note S1. Publisher supplement.
- [3] SUMMIT manuscript and Supplementary Information, local working version, Methods “Estimation of genome-wide LD scores” and Supplementary Note “Effect of reference-LD mismatch on SUMMIT estimates” (accessed 13 August 2026).
- [4] GENIE source, commit 22bee9c, covariate inverse and projection. `ge_flexible.cpp`, lines 2303–2362.
- [5] GENIE source, commit 22bee9c, genotype mean and HWE scale. `ge_flexible.cpp`, lines 654–665.
- [6] GENIE source, commit 22bee9c, additive trace initialization. `ge_flexible.cpp`, lines 1655–1674.
- [7] GENIE source, commit 22bee9c, projected phenotype normalization. `ge_flexible.cpp`, lines 2331–2362.